


# Introduction to Machine Learning

Badri Narayan Subudhi  
IIT Jammu

# Machine Learning




- Machine learning is a branch of AI focused on building computer systems that learn from data. The breadth of ML techniques enables software applications to improve their performance over time.
- Traditional machine learning combines data with statistical tools to predict outputs, yielding actionable insights.

# What is Machine Learning?


- Pattern Recognition is the study of how machines can:
  - observe the environment,
  - learn to distinguish patterns of interest,
  - make sound and reasonable decisions about the categories of the patterns.
- What is a *pattern*?
- What kinds of *category* we have?

# What is a pattern?

- A **pattern** is an abstraction, represented by a set of measurements describing a “physical” object.
- Many types of patterns exist:
  - visual, temporal, sonic, logical, ...




John Smith



# What is a Pattern Class (or category)?

- Is a set of patterns sharing common attributes
- A collection of “similar”, not necessarily identical, objects



# What is **Features** ?

- Features are properties of an object:
  - Ideally representative of a specific type (i.e. class) of objects
  - Perceptual relevant



# Good features vs. bad features

- Ideally, for a given group of patterns coming from the same class, feature values should all be similar
- For patterns coming from different classes, the feature values should be different.



# Applications

## Recognition (OCR)



Printed texts: reading machines for blind people, digitalization of text documents.

## Biometrics



- Face recognition, verification, retrieval.
- Finger prints recognition.
- Speech recognition.

## Diagnostic systems



- Medical diagnosis: X-Ray, ECG analysis.
- Machine diagnostics, waster detection.

## Target Tracking



Object detection and tracking from videos



# Classification in Statistical ML

- A **feature extractor** is a program that inputs the data (image) and extracts features that can be used in classification.
- A **classifier** is a program that inputs the feature vector and assigns it to one of a set of designated classes or to the “reject” class.

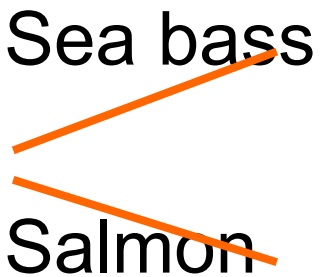
***With what kinds of classes do you work?***


# ML Process

- **Data acquisition and sensing:**
  - Measurements of physical variables.
  - Important issues: bandwidth, resolution , etc.
- **Pre-processing:**
  - Feature selection.
  - Isolation of patterns of interest from the background.
- **Feature extraction:**
  - Finding a new representation in terms of features.
- **Classification**
  - Using features and learned models to assign a pattern to a category.

# An Example

“Sorting incoming Fish on a conveyor according to species using optical sensing”

Species   
Sea bass  
~~Salmon~~




- Problem Analysis

- Set up a camera and take some sample images to extract features
  - Length
  - Lightness
  - Width
  - Number and shape of fins
  - Position of the mouth, etc...


This is the set of all suggested features to explore for use in our classifier!

- Preprocessing
  - Use a segmentation operation to isolate fishes from one another and from the background
- Information from a single fish is sent to a **feature extractor** whose purpose is to represent the data by measuring certain features
- The features are passed to a **classifier**



- Classification

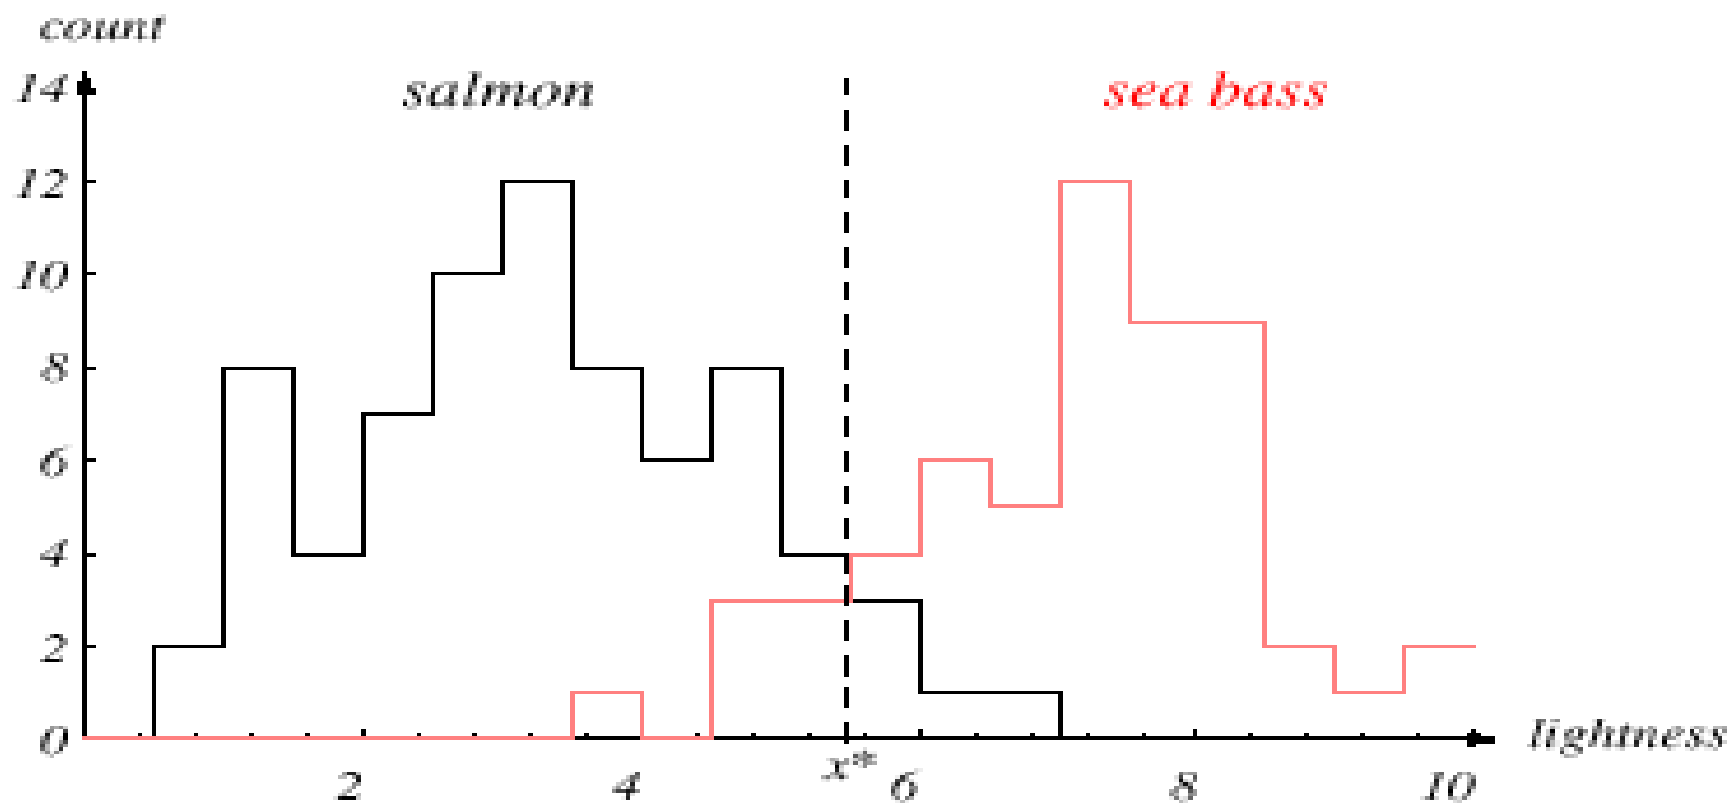
Select the length of the fish as a possible feature for discrimination






The **length** is a poor feature alone!

Select the **lightness** as a possible feature.




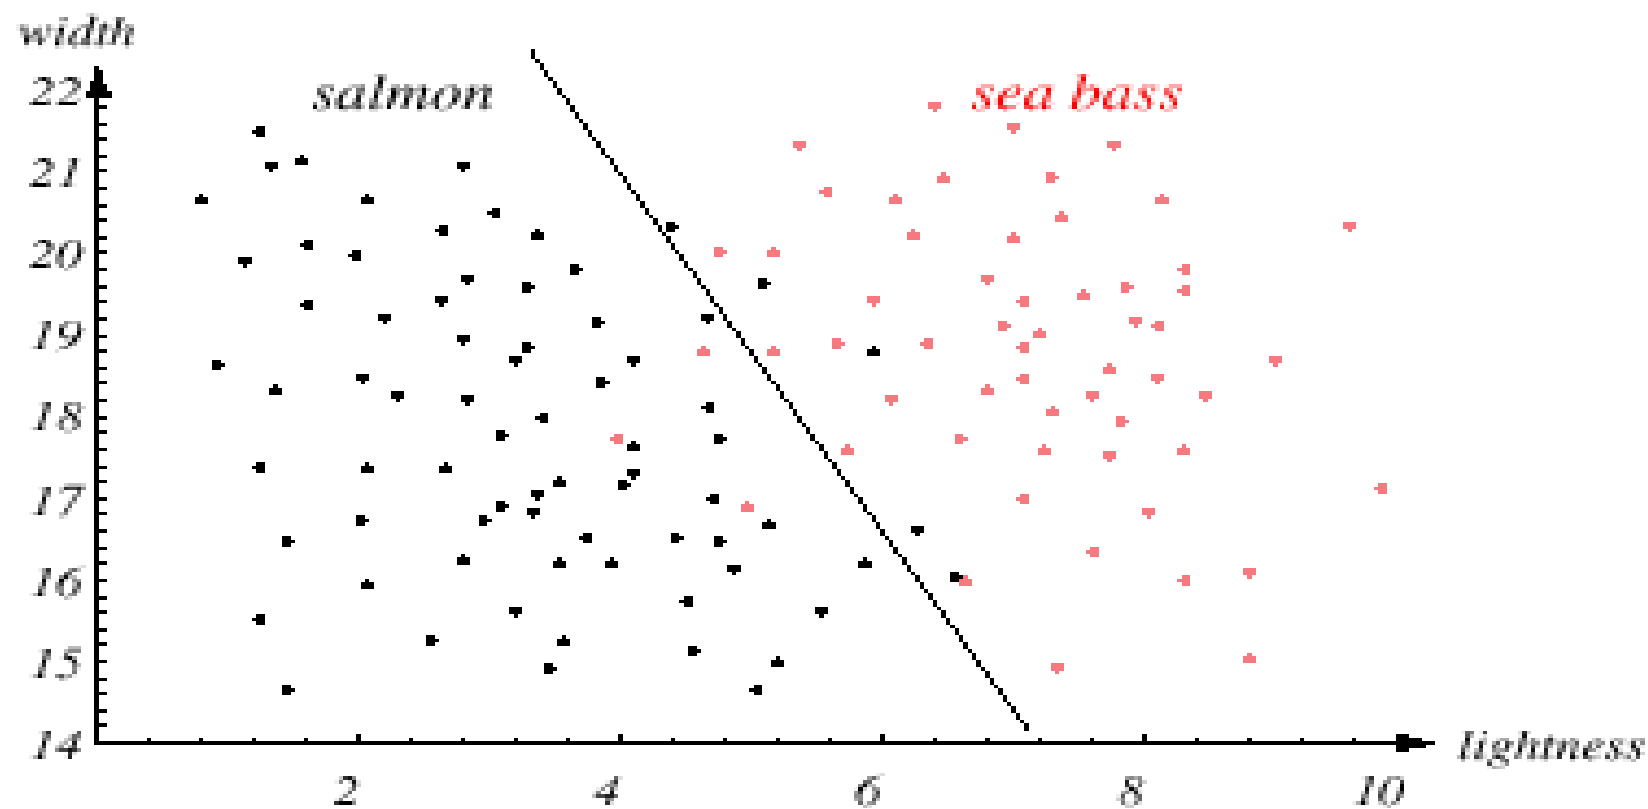
- Threshold decision boundary and cost relationship
  - Move our decision boundary toward smaller values of lightness in order to minimize the cost (reduce the number of sea bass that are classified salmon!)



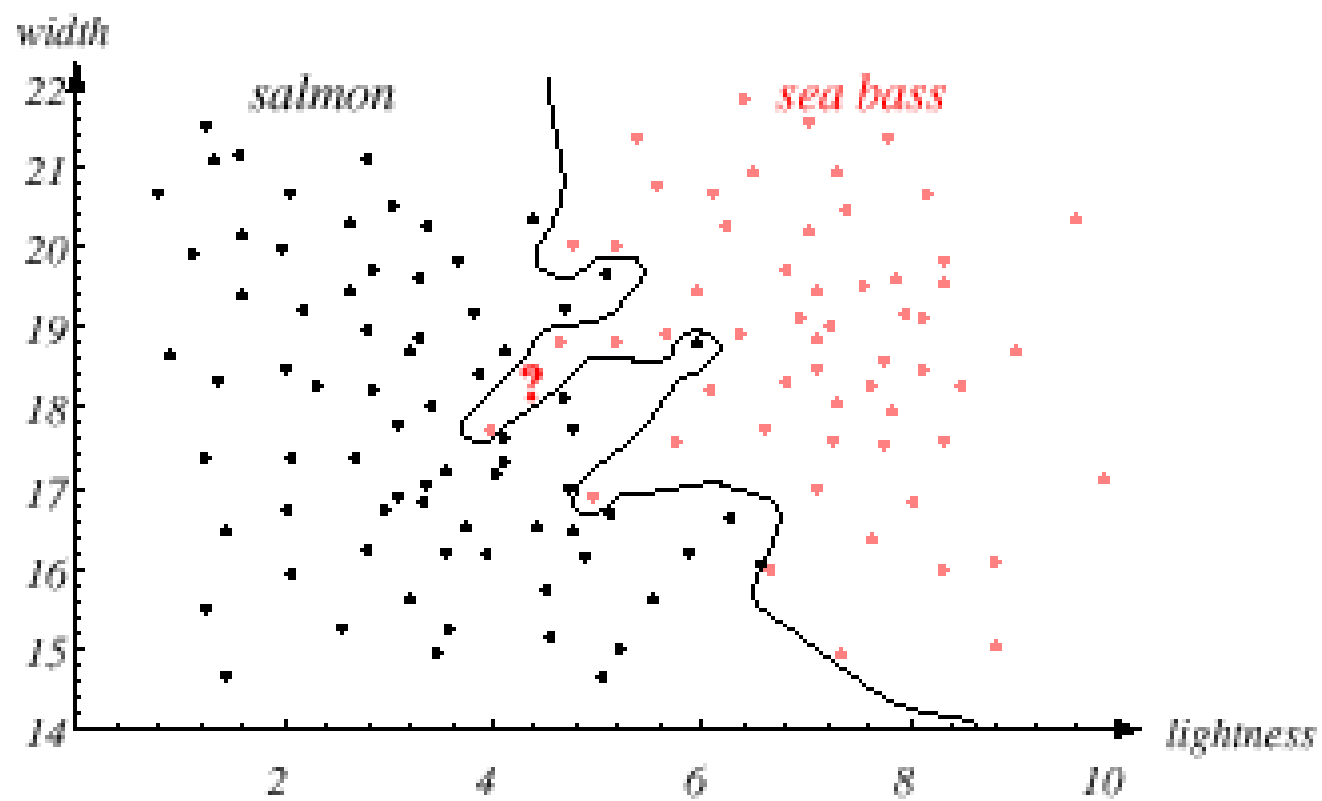
Task of decision theory

- Adopt the lightness and add the width of the fish






- We might add other features that are not correlated with the ones we already have. A precaution should be taken not to reduce the performance by adding such “noisy features”
- Ideally, the best decision boundary should be the one which provides an optimal performance such as in the following figure:




- However, our satisfaction is premature because the central aim of designing a classifier is to correctly classify novel input



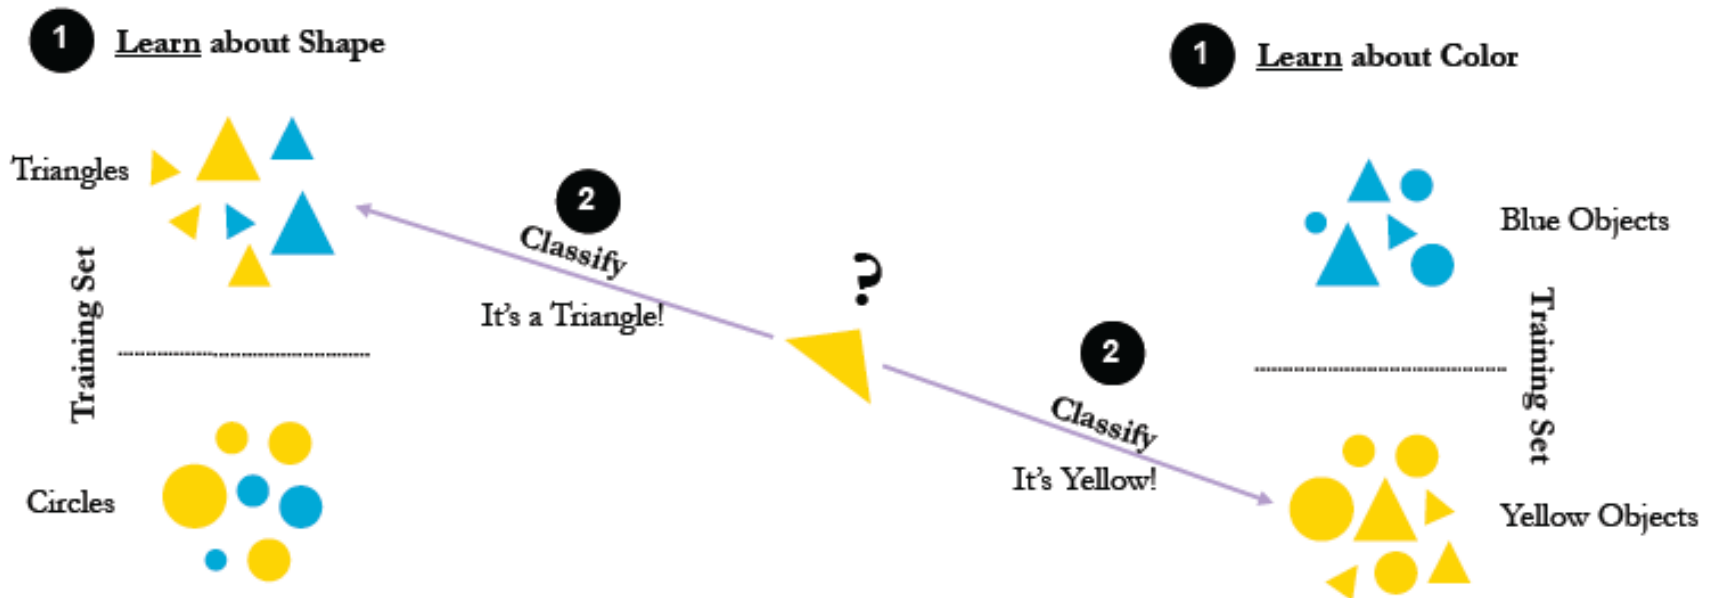
Issue of generalization!






# Supervised Training/Learning

– a “teacher” provides labeled training sets, used to train a classifier







# Classifier

A classifier partitions sample space  $X$  into **class-labeled regions** such that

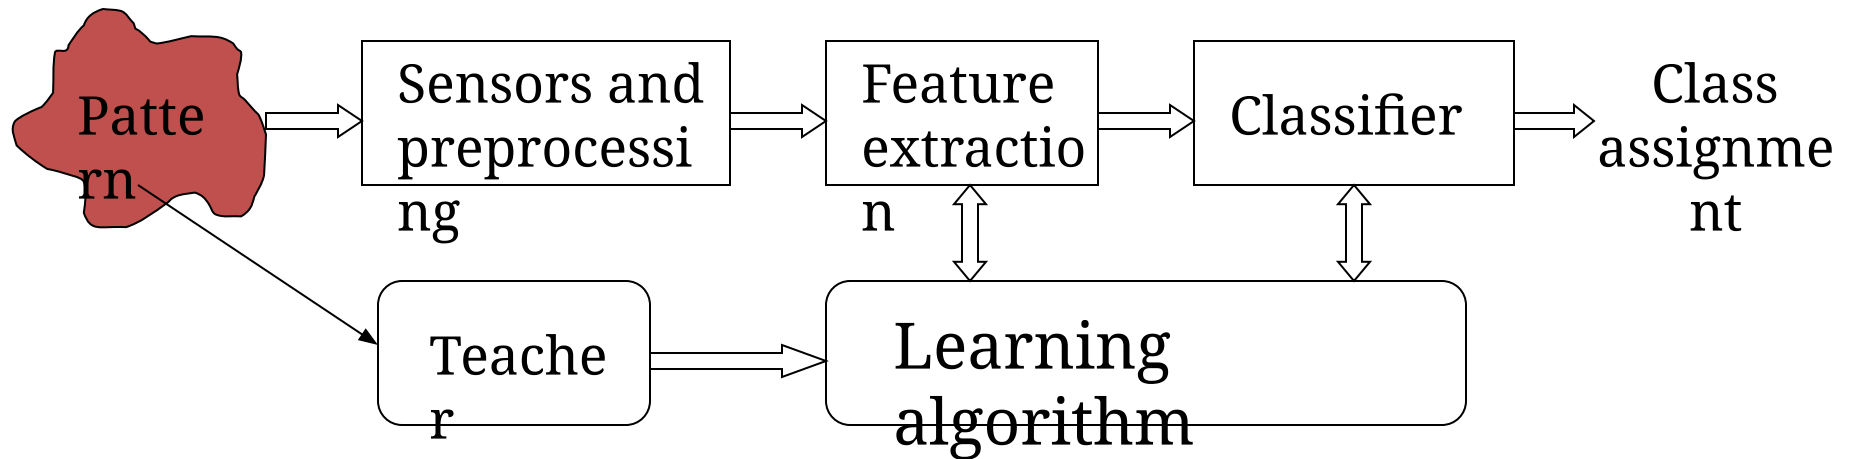
$$X = X_1 \cup X_2 \cup \dots \cup X_{|Y|} \quad \text{and} \quad X_i \cap X_j = \{\emptyset\}$$



The classification consists of determining to which region a feature vector  $\mathbf{x}$  belongs to.


Borders between **decision boundaries** are called decision regions.

# Components of classifier system



- **Sensors and preprocessing.**
- **A feature extraction** aims to create discriminative features good for classification.
- **A classifier.**
- **A teacher** provides information about hidden state -- supervised learning.
- **A learning algorithm** sets PR from training examples.

# Basic concepts



## Feature vector $\mathbf{x} \in X$

- A vector of observations (measurements).

-  $\mathbf{x}$  is a point in feature space  $X$ .

## Hidden state $y \in Y$

- Cannot be directly measured.

- Patterns with equal hidden state belong to the same class.

## Task

- To design a classifier (decision rule)  $q: X \rightarrow Y$


which decides about a hidden state based on an observation.

# Training Samples

When a few labeled patterns can be collected by experts, Supervised learning can be opted instead of the unsupervised approach to make full utilization of labeled patterns.

In a dataset a training set can be implemented to build up a model.

# Example



Task: American-Indian recognition.


The set of hidden states  $\mathcal{Y} = \{A, I\}$

The feature space  $\mathcal{X} = \mathbb{R}^2$

Linear classifier:

$$q(\mathbf{x}) = \begin{cases} A & \text{if } (\mathbf{w} \cdot \mathbf{x}) + b \geq 0 \\ I & \text{if } (\mathbf{w} \cdot \mathbf{x}) + b < 0 \end{cases}$$


Training examples






# Overfitting and underfitting


**Problem:** how rich class of classifications  $q(\mathbf{x};\theta)$  to use.



underfitting



good fit

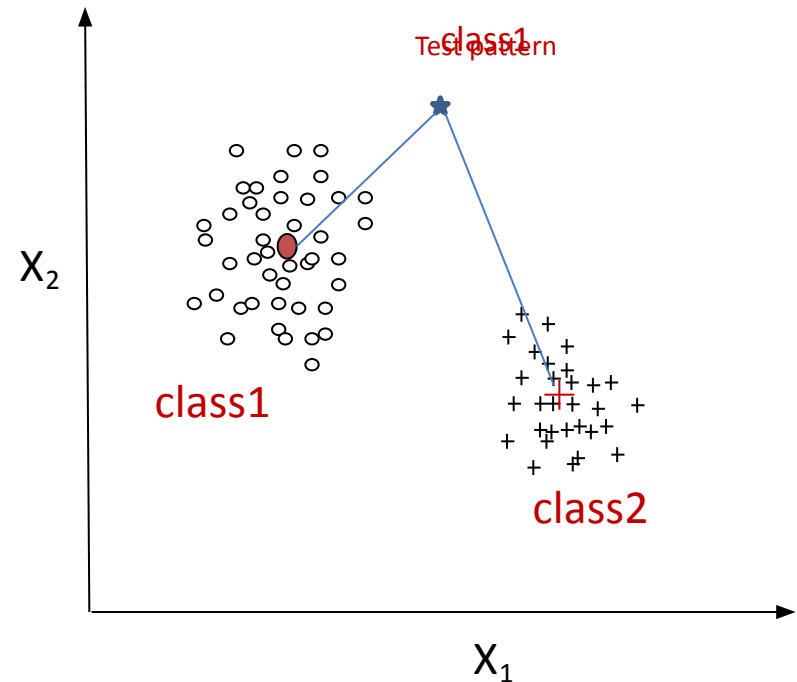


overfitting

**Problem of generalization:** a small empirical risk  $R_{\text{emp}}$  does not imply small true expected risk  $R$ .

# Minimum Distance Classifier

- Each class is represented by its mean vector
- Training is performed by calculating the mean of the feature vectors of each class
- New patterns are classified by finding the closest mean vector
- The boundary is the perpendicular bisector of the line joining the mean points.



# Minimum Distance Classifier

$$m_j = \frac{1}{N_j} \sum_{X \in \omega_j} X \quad j = 1, 2, \dots, M$$

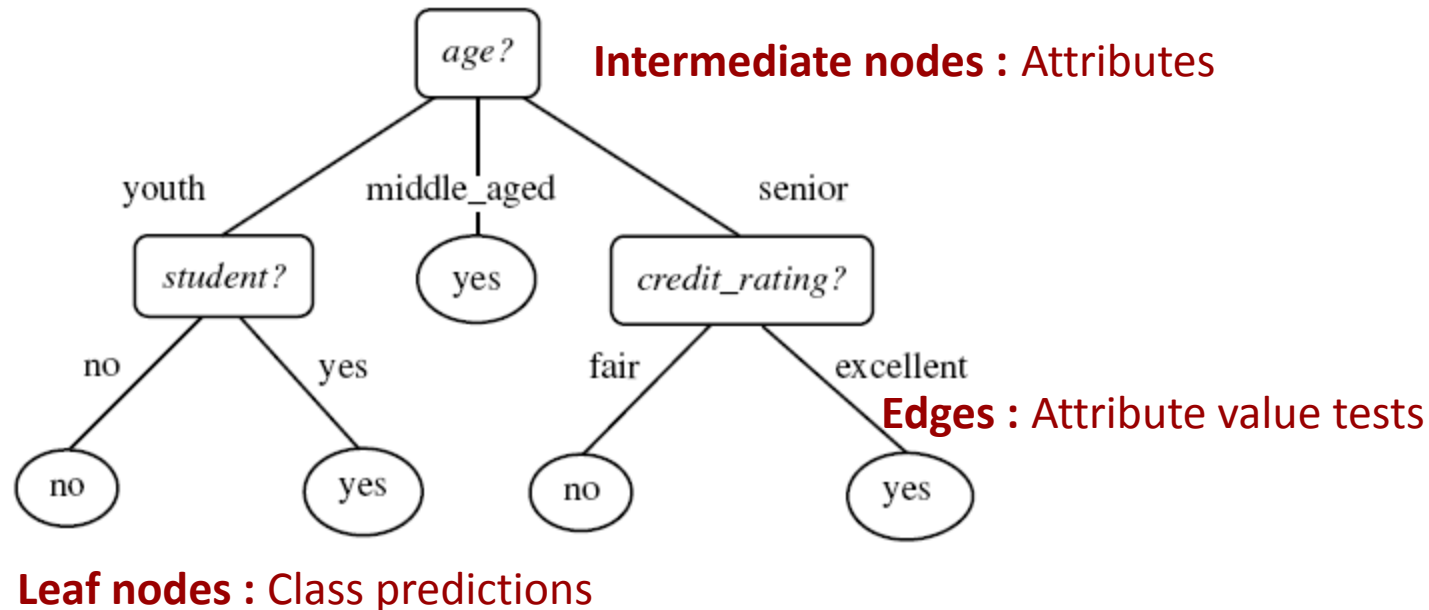
$N_j$  = number of pixels from class  $\omega_j$

$$D_j(X) = \|X - m_j\| \quad (\text{Euclidean distance})$$

Assign  $X$  to  $\omega_i$  if  $D_i(X) < D_j(X), \quad j = 1, 2, \dots, M; j \neq i$


# Decision Trees

# Example tree



**Example algorithms:** ID3, C4.5, SPRINT, CART

# Decision Tree schematic



# Decision Tree Issues

## How to determine the attribute for split?

Alternatives:

1. Information Gain

$$\text{Gain}(A, S) = \text{Entropy}(S) - \sum (S_j/S) * \text{Entropy}(S_j)$$

Other options:

Gain ratio, etc.

# Lazy learners



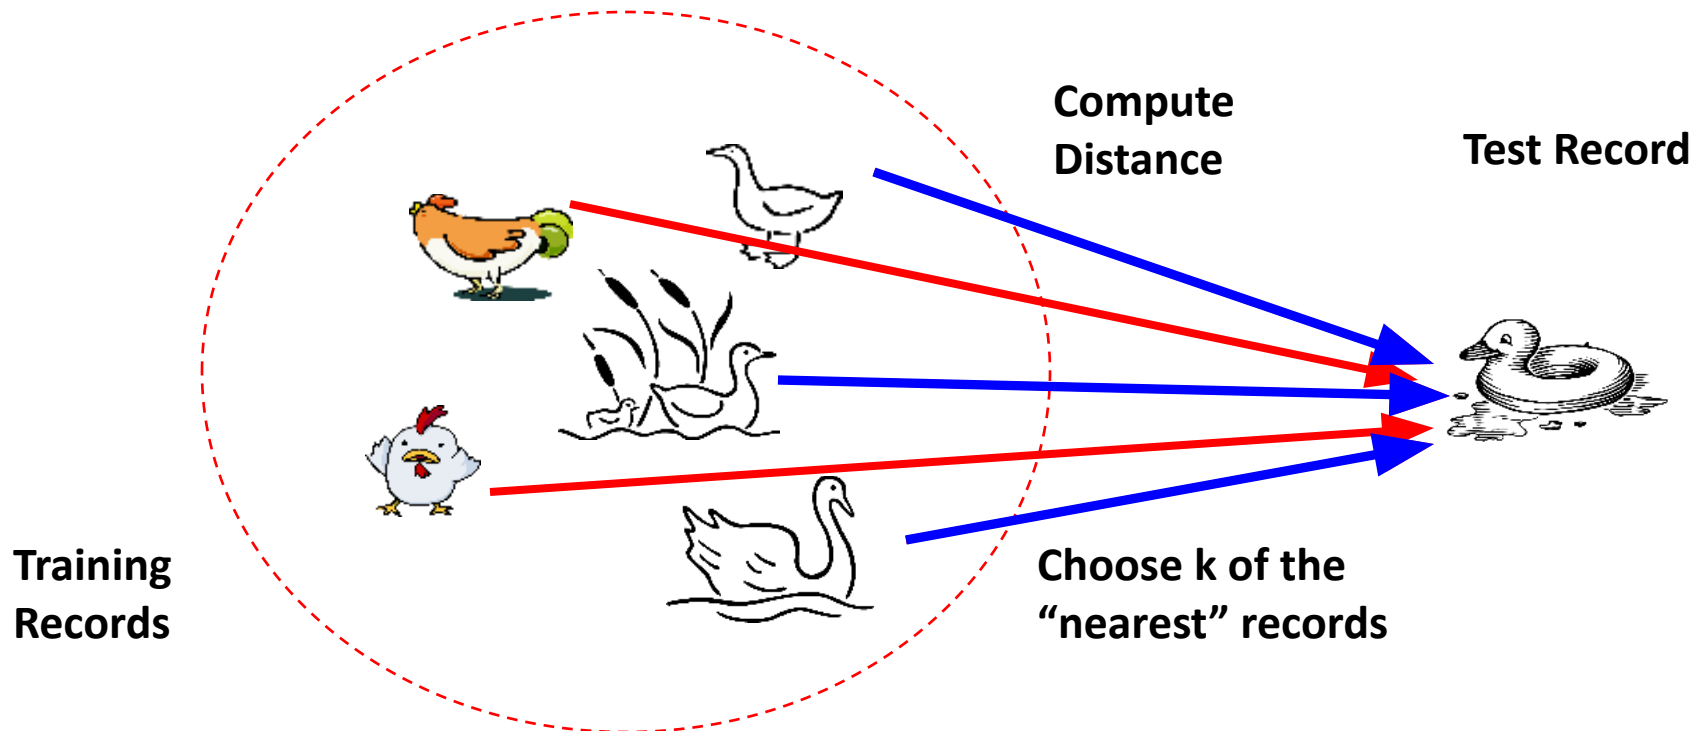
# Lazy learners

- ‘**Lazy**’: Do not create a model of the training instances in advance
- When an instance arrives for testing, runs the algorithm to get the class prediction
- Example, K – nearest neighbour classifier (K – NN classifier)

“One is known by the company one keeps”

# Pattern-Based Classification: Nearest Neighbor Classifiers

- Basic idea:
  - If it walks like a duck, quacks like a duck, then it's probably a duck




# K-NN classifier schematic

For a test instance,

- 1) Calculate distances from training pts.
- 2) Find K-nearest neighbours (say, K = 3)
- 3) Assign class label based on majority

$$\text{dist}(X_1, X_2) = \sqrt{\sum_{i=1}^n (x_{1i} - x_{2i})^2}.$$

$$v' = \frac{v - \min_A}{\max_A - \min_A},$$



# K-NN classifier Issues

**How to determine distances between values of categorical attributes?**

Alternatives:

1. Boolean distance (1 if same, 0 if different)
2. Differential grading (e.g. weather – ‘drizzling’ and ‘rainy’ are closer than ‘rainy’ and ‘sunny’ )

# Bayesian Classification

We will pose the classification problem in  
probabilistic terms

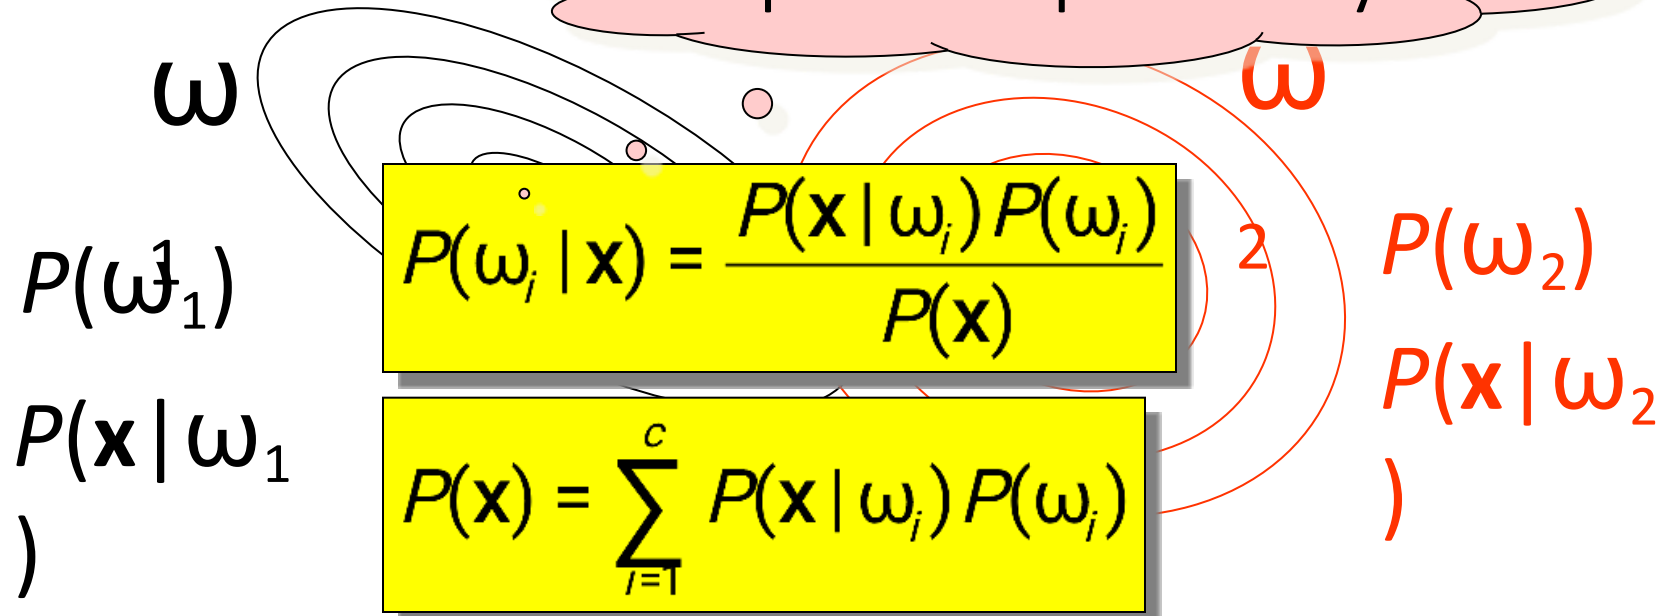
Create models for how features are distributed  
for objects of different classes

We will use probability calculus to make  
classification decisions

# The States of Nature


$P(\mathbf{x} | \omega_i)$ : class-conditional probability

A posterior probability



# Lets Look at Just One Feature


- Each object can be associated with multiple features
- We will look at the case of just one feature for now



We are going to define two key concepts....

# The First Key Concept

Features for each class drawn from **class-conditional probability distributions (CCPD)**



Our first goal will be to *model* these distributions



# The Second Key Concept

We model **prior probabilities** to quantify the expected *a priori* chance of seeing a class

$$P(\text{Class2}) \quad \& \quad P(\text{Class1})$$

# But How Do We Classify?

- So we have priors defining the *a priori* probability of a class

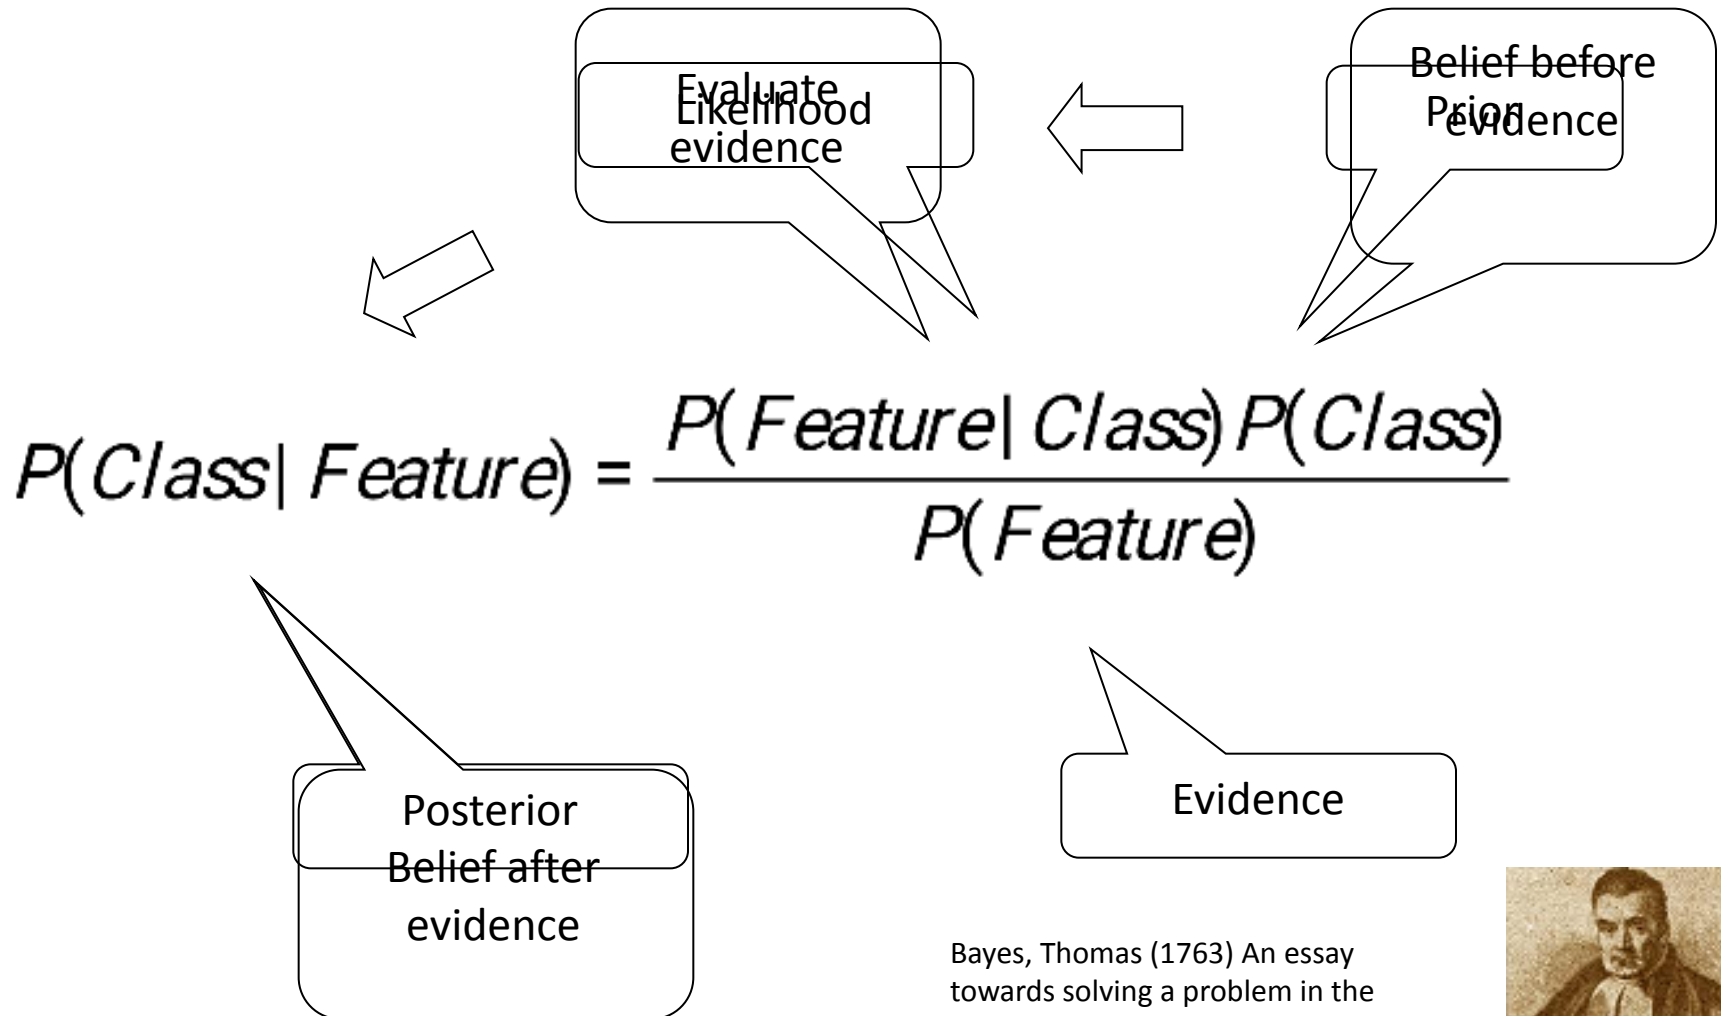
$$P(\text{Class1}), P(\text{Class2})$$

- We also have models for the probability of a feature given each class

$$P(X|\text{Class1}), P(X|\text{Class2})$$

*But we want the probability of the class given a feature  
How do we get  $P(\text{Class1}|X)$ ?*

# Bayes Rule



Bayes, Thomas (1763) An essay towards solving a problem in the doctrine of chances. Philosophical Transactions of the Royal Society of London, 53:370-418



# Bayes Decision Rule

If we observe an object with feature  $X$ , how do decide if the object is from Class 1?

The **Bayes Decision Rule** is simply choose Class1 if:

This is the same number on both sides!

# Discriminant Function

We can create a convenient representation of the Bayes  
Decision Rule

$$P(X | \text{Class1}) P(\text{Class1}) > P(X | \text{Class2}) P(\text{Class2})$$

*If  $G(X) > 0$ , we classify as Class 1*

# Stepping back

What do we have so far?

We have defined the two components, **class-conditional distributions** and **priors**

$$P(X | \text{Class1}), P(X | \text{Class2}) \quad P(\text{Class1}), P(\text{Class2})$$

We have used Bayes Rule to create **a discriminant function for classification** from these components

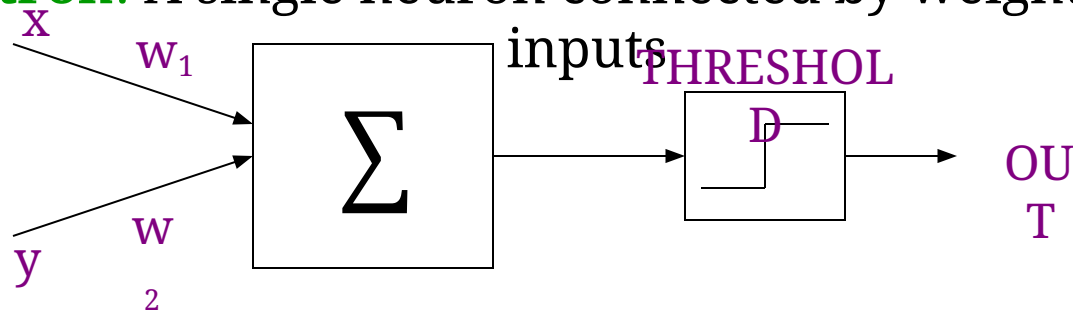
$$G(X) = \log \frac{P(X | \text{Class1})}{P(X | \text{Class2})} \frac{P(\text{Class1})}{P(\text{Class2})} > 0$$

Given a new feature, X, we plug it into this equation...


...and if  $G(X) > 0$  we classify as **Class1**

# Neural Classifier

**Perceptron:** A single neuron connected by weights to a set of



- Let  $x$  &  $y$  be two inputs and  $w_1, w_2$  be the weights.
- If  $w_1x + w_2y > \theta$  then the output is **1** else **0**, where  $\theta$  = threshold
  - $w_1x + w_2y = \theta \rightarrow$  separating line



## Learning rule

**Learning:** Present a set of input patterns, adjust the weights until the desired output occurs for each of them.

$$w_i(t+1) = w_i(t) + \Delta_i;$$


$$\Delta_i = \eta \delta x_i;$$

$$\delta = T - A \text{ (i.e., target - actual).}$$

- ❖ If the sets of patterns are linearly separable, the single layer perceptron algorithm is guaranteed to find a separating hyperplane in a finite number of steps.




# Change of weights





# Cascading layers

*Two layers :*



*Three layers:*



*Multi-layer network*

# Application for Change Detection




Image of date-1







image of date-2




Difference image



Min-distance



KNN






Bayes Classifier



Neural network

## Video Tracking

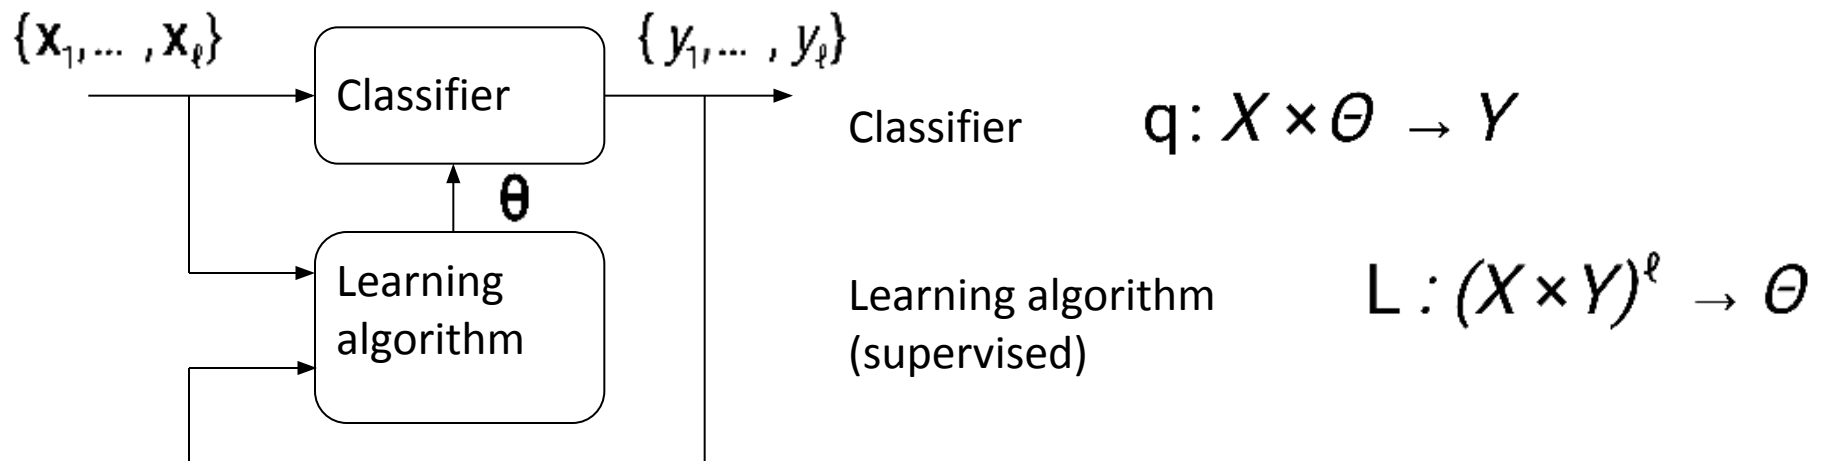


# Unsupervised learning/ clustering

Input: Data  $\{\mathbf{x}_1, \dots, \mathbf{x}_\ell\}$  without information about the hidden state.

Clustering: goal is to find clusters of data sharing similar properties.

A broad class of **unsupervised learning algorithms**:

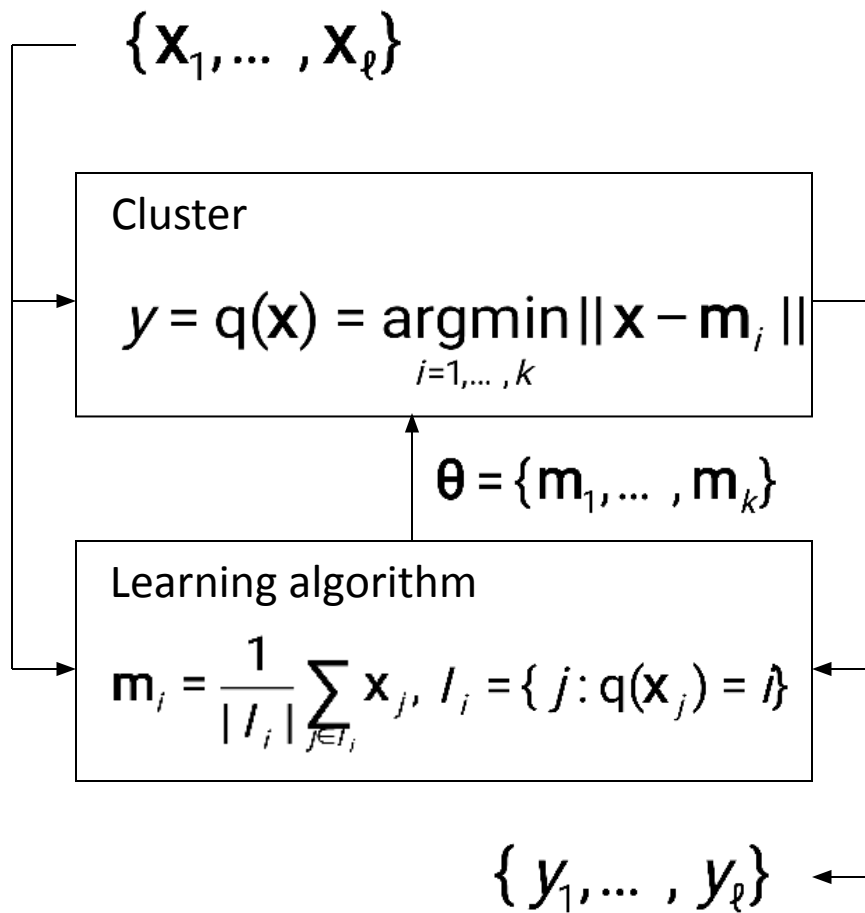


# Good Clustering

- A good clustering method will produce high quality clusters with
  - high intra-class similarity
  - low inter-class similarity
- The quality of a clustering result depends on both the similarity measure used by the method and its implementation.
- The quality of a clustering method is also measured by its ability to discover some or all of the hidden patterns.


# Example of unsupervised learning algorithm

## k-Means clustering:



Goal is to minimize

$$\sum_{i=1}^{\ell} \|\mathbf{x}_i - \mathbf{m}_{q(\mathbf{x}_i)}\|^2$$





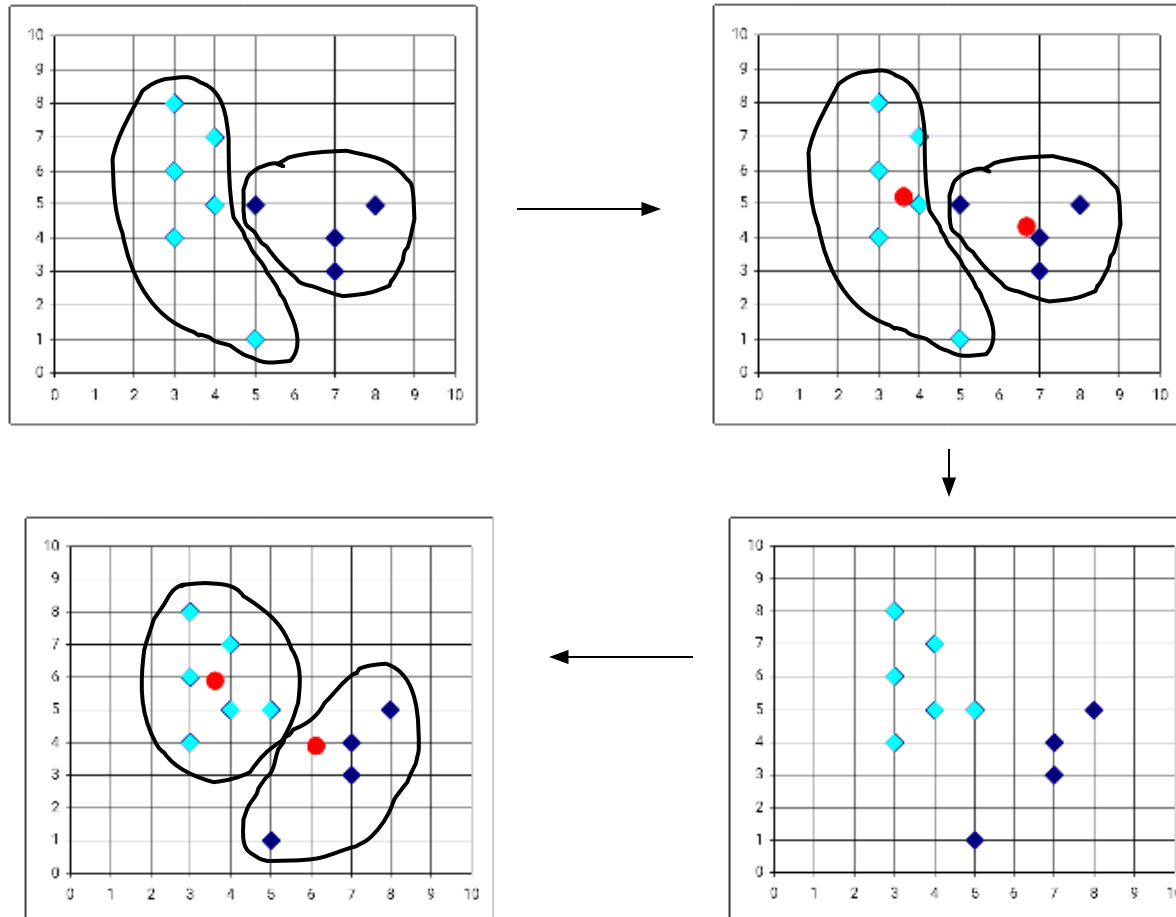
# K-means clustering algorithm

Given  $k$ , the *k-means* algorithm is implemented in 4 steps:

- i. Partition objects into  $k$  nonempty subsets
- ii. Compute seed points as the centroids of the clusters of the current partition.
- iii. The centroid is the center (mean point) of the cluster.
- iv. Assign each object to the cluster with the nearest seed point.
- v. Go back to Step iii, stop when no more new assignment.

# K-means clustering algorithm

- Example



# References

## **Books**

Theodoridis and Koutroumbas, Pattern Recognition. Academic Press, San Diego, 2003.

Duda and Hart: Pattern Classification and Scene Analysis. J. Wiley & Sons, New York, 1982. (2nd edition 2000).

Fukunaga: Introduction to Statistical Pattern Recognition. Academic Press, 1990.

Bishop: Neural Networks for Pattern Recognition. Claredon Press, Oxford, 1997.

# Thank you for your attention!



"But before we move on, allow me to belabor the point even further..."

## Questions?

